

Line and point detection

1 Feature point detection and matching

Exercise 1. Write a function `corners(image: np.ndarray, w: int, kappa: float)` that implements the Harris corner detector. Your function should:

- Apply a $w \times w$ Gaussian filter to the image I to remove noise.
- Compute the gradients I_x , I_y of the image in x and y directions.
- Compute the Harris operator at each pixel (x, y)

$$K(x, y) = \frac{\det(H(x, y))}{\text{trace}(H(x, y))} \text{ where } H(x, y) = \sum_{(u, v) \in W(x, y)} w(u, v) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}_{(x+u, y+v)}$$

where $w(u, v)$ is the Gaussian window function and $W(x, y)$ is a $w \times w$ neighborhood of (x, y) .

- Find the set of pixels $H = \{(x, y) : K(x, y) > \kappa\}$ such that the response of the Harris operator is above a threshold κ .
- Apply non-maximum-suppression in a $w \times w$ pixel neighborhood of each $(x, y) \in H$. Store the pixel coordinates of the resulting keypoints.

Write a script that tests your code on images “ames1.JPG” for different values of the window size w and the Harris corner threshold κ . For the best choice, plot the results of each step, as well as the Harris corners before and after non-maximum-suppression.

Exercise 2. Write a function `features(image: np.ndarray, keypoints: np.ndarray)` that extracts a descriptor centered around each keypoint in $\text{keypoints} \in \mathbb{R}^{2 \times P}$. Use a simple vectorized image patch of 9×9 pixels as your descriptor so that $\text{descriptors} \in \mathbb{R}^{81 \times P}$.

Exercise 3. Write a function `matching(descriptors1: np.ndarray, descriptors2: np.ndarray, tau: float)` that matches two sets of feature descriptors using the sum of squared differences (SSD) as a matching score. Your function should return for each descriptor in image 1 the index of the best match in image 2 provided that the SSD score is above a threshold τ . If a descriptor in image 1 has no match, then the corresponding entry of matches should be set to -1 . Write a script that tests your code on image pair: “ames1.JPG” and “ames2.JPG” and plots the resulting matches. Plot also the ROC curve and use it to select the threshold.

2 Line detector

Exercise 4. Write a script that does the following.

- Apply the **Canny edge detector** to images: “ames1.JPG”.
- Apply **Hough transform** to the detected edges. Plot the lines estimated by Hough transform superimposed on the image.
- Show the original image, edge map (after step 1), and final image (after step 2).